

Notes on transient mortality shocks in low-mortality populations

Joshua R. Goldstein Ronald D. Lee

May 17, 2020

Abstract

Here we present preliminary results for the modeling of transient mortality shocks based on Lee-Carter estimates of the evolution of mortality over time. The model we fit includes random terms for the evolution of the underlying trend, which is modeled as a random walk with deterministic drift, and for transient shocks, which we interpret as single-year events due to weather or annual infectious diseases like the flu. These preliminary models only work well for some countries. Nonetheless it is still possible to see that the transient shocks of neighbors are correlated, suggesting that they are picking up real variation in mortality, rather than measurement error.

1 Overview

The canonical time-series model used for the Lee-Carter model is the random walk with drift. The logic behind this is that there is a steady march of progress toward lower mortality which varies in its pace. The recent three-year run of declining mortality in the United States – and the current pandemic – suggest that there are short-term factors other than technological progress and changes in population health that may be important for mortality in any given year.

Modeling these short-term transient shocks is interesting in its own right, because of what it reveals about the nature of population mortality levels and changes. It may also be useful for understanding the long-run implications

of mortality reversals like that from the U.S. opioid crisis or the worldwide coronavirus pandemic.

Here we present preliminary results for the modeling of transient mortality shocks based on Lee-Carter estimates of the evolution of mortality over time. The model we fit includes random terms for the evolution of the underlying trend, which is modeled as a random walk with deterministic drift, and for transient shocks, which we interpret as single-year events due to weather or annual infectious diseases like the flu.

These preliminary models only work well for some countries. Nonetheless it is still possible to see that the transient shocks of neighbors are correlated, suggesting that they are picking up real variation in mortality, rather than measurement error.

2 Modeling

We first estimate the Lee-Carter model, adjusting k_t to match e_0 . This model reduces the logarithm of the full set of age-period mortality rates to an average age-schedule a_x and a time index k_t that drives age-specific changes b_x . The model has the form

$$\log M_{x,t} = a_x + b_x k_t.$$

Standard methods are available for estimating the parameters.

The usual time series model for forecasting is the random walk with drift

$$k_t = k_{t-1} + d + \epsilon_t.$$

Here we add another layer to the estimation of the time series, decomposing the observed k_t into a latent k_t that evolves as a random walk with drift and an annual transitory component n_t . In state-space modeling n_t is sometimes called “observation error” or “noise”. We are conceiving of it not so much as error but as a transitory perturbation – for example due to weather or the severity of the annual flu or to another kind of contagious disease such as COVID-19.

The model has the form

$$k_t^{observed} = k_t^{latent} + n_t \tag{1}$$

$$k_t^{latent} = k_{t-1}^{latent} + d + \epsilon_t \quad (2)$$

The model has two features.

- The observed value is the latent value plus “noise” n_t , which is assumed to be normally distributed with constant variance and independent over t .
- The latent value evolves as a random walk with drift, with a fixed (deterministic) value of d .

Note it is also possible to add an additional random layer to this model by making d itself a stochastic evolving term. The standard form of this “random trend” model is to let d evolve as a random walk, $d_t = d_{t-1} + \eta_{t-1}$. We don’t consider this model for now.

3 Results

In Figure 1 we show the observed and latent k_t for 9 countries in the Human Mortality Database. In some countries, such as the United States, Russia, and West Germany, the latent level of mortality is indistinguishable from the observed level. In the others, latent value is clearly a smoother version of the observed k_t . The smoothing is particularly evident in France, Spain, and Sweden. Italy and Japan are intermediate cases.

The contrast between countries can be seen more clearly in Figure 2, which shows the same estimates as in Figure 1 but only since 1990.

Figure 3 shows the estimated values of n_t (the difference between observed and latent k_t). These can be interpreted as the transitory shocks (and/or measurement error). The plots are done on the same scale. The large values are associated with countries in which there is a lot of smoothing.

Finally, 3 shows how the “noise” n_t is correlated for neighboring countries. (Note: The scale differs by plot). We can see very high correlations between Italy and France (0.76) and very low correlations between the UK and Japan (0.13).

Interestingly, we see that Italy, Spain, and France are all closely related. So are the United States and Canada. Sweden in the United States are also similar.

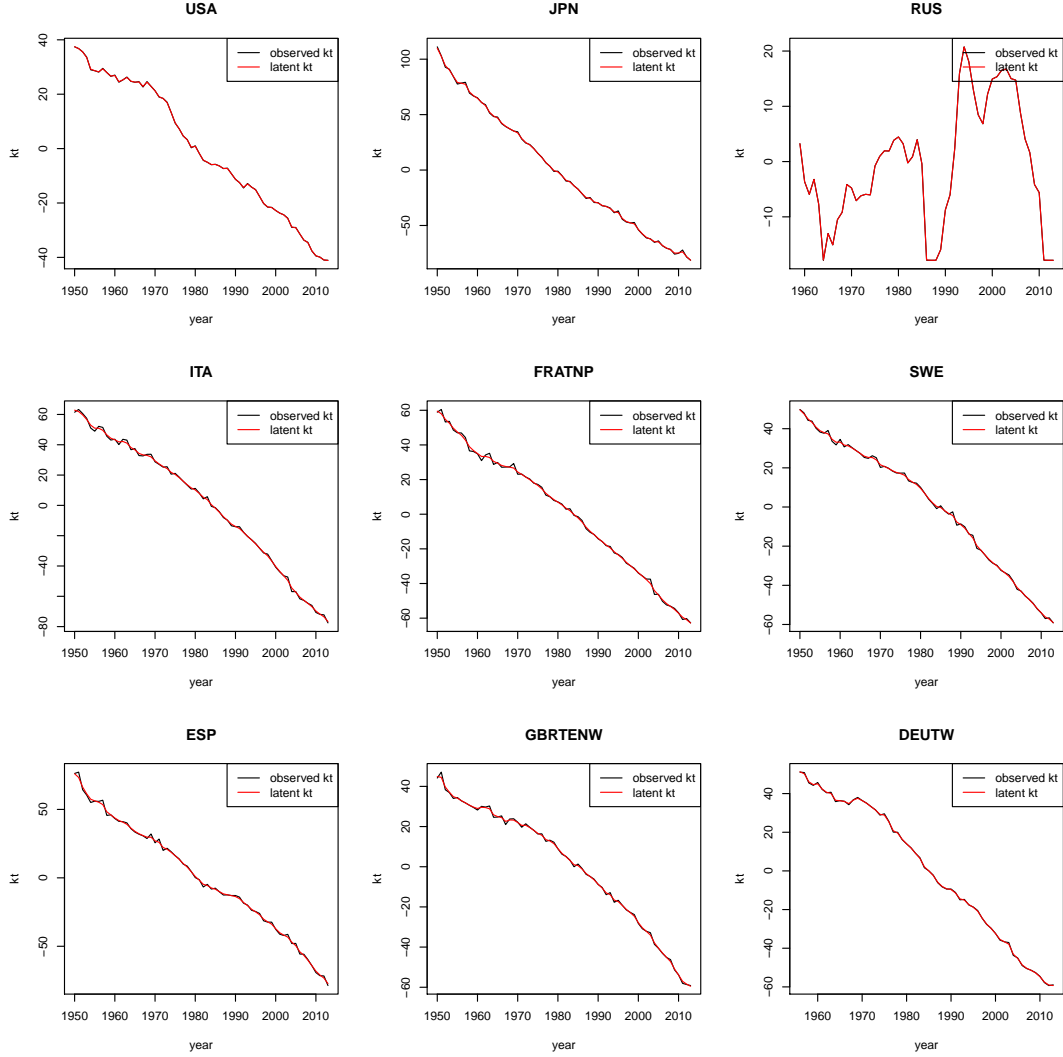


Figure 1: Estimates of observed and latent values of k_t in Lee-Carter model using MARSS package based on HMD data. Zoom in to see detail. Take home: In some countries (e.g., Sweden, France, Spain, and the UK), the estimation gives a much smoother latent value. In the USA this model doesn't produce much of interest.

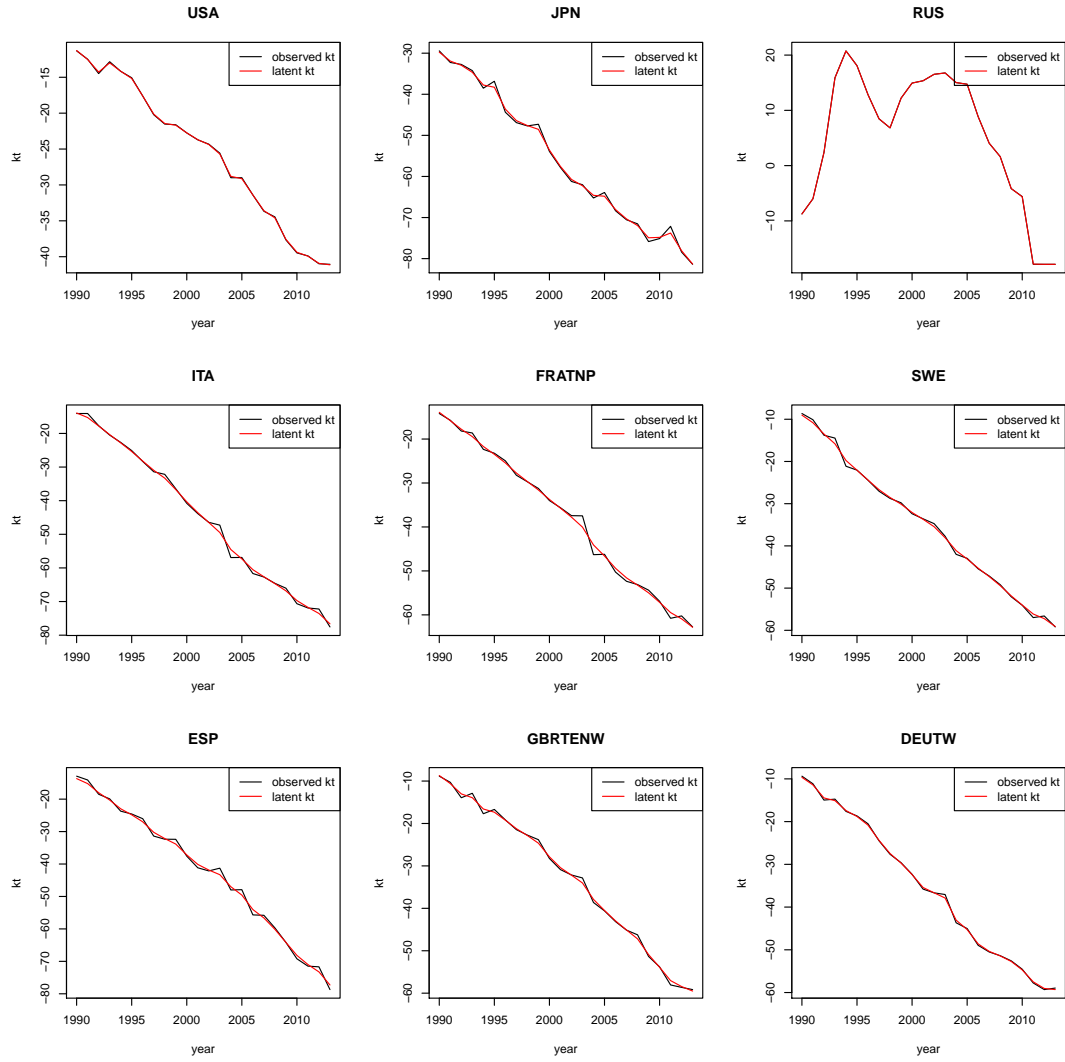


Figure 2: Zoomed in version of Figure 1, showing results since 1990.

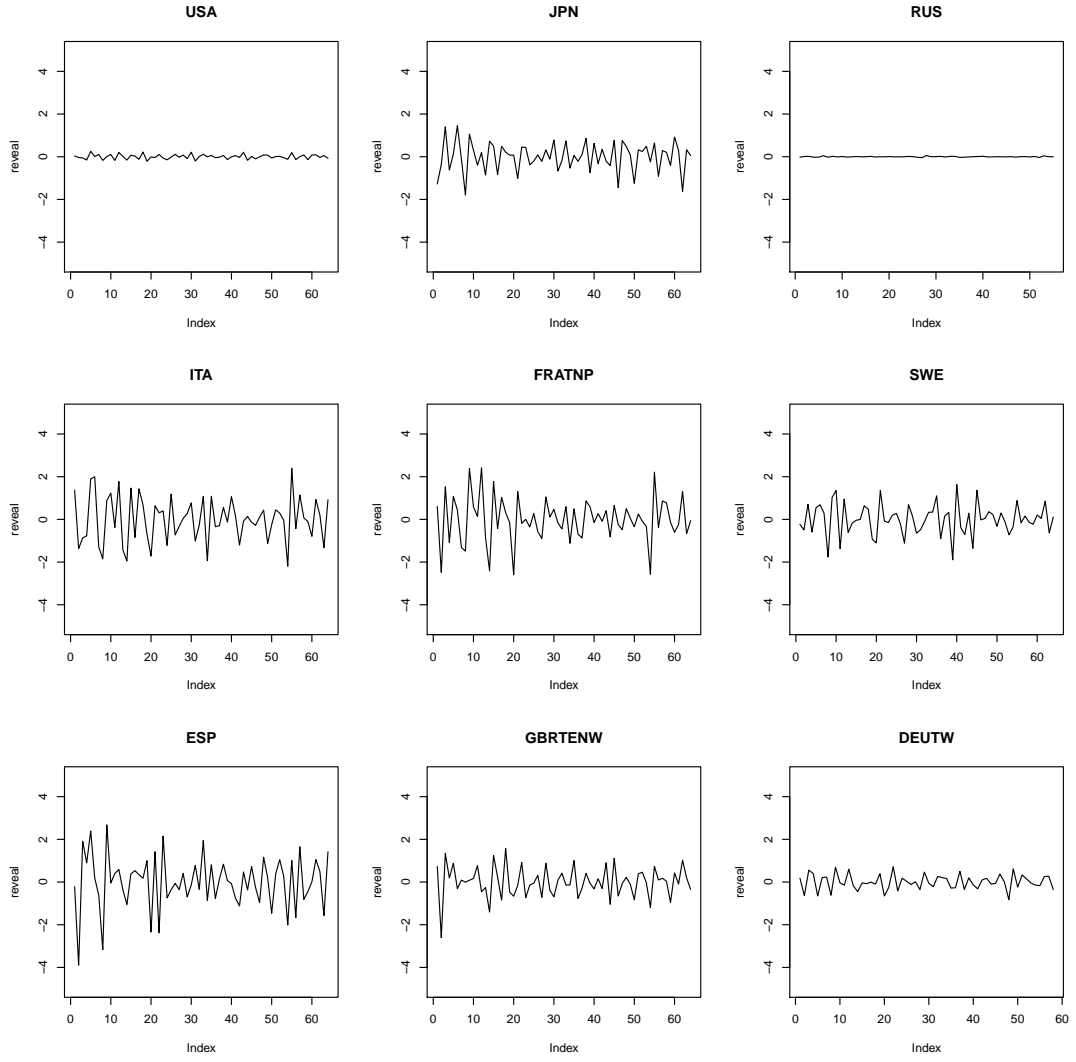


Figure 3: Estimated shocks $n_t = k_t^{latent} - k_t^{observed}$, shown on a common scale.




Figure 4: Correlations across countries in the time series of $n_t(\text{country})$. Note strong correlations between Spain, France, and Italy – and also the USA and Canada.

4 Discussion

Our estimation of transitory mortality is not complete. But a few preliminary conclusions can be reached

1. In some countries, the observed k_t is noticeably more volatile than the latent k_t . In these cases, it might be preferable to begin forecasts at the most recent value of the latent k_t . One might also want to use the estimated variances of ϵ_t and n_t to estimate forecast uncertainty.
2. The “noise” we are estimating is correlated for neighboring countries. This makes it appear that there are weather and infectious diseases “shocks” that are shared across countries. It also makes it unlikely that measurement issues are behind the shocks.
3. This simple model does not seem to apply to United States, where the latent and observed k_t are indistinguishable. This is consistent with Lawrence Carter’s paper, showing that structural time series forecasts for the US are nearly the same as ARIMA

5 Future directions

Future steps to be taken could include

1. Check on convergence, measures of model fit and so forth for all of these countries. (The estimates shown should be considered preliminary.)
2. Check to see if making the drift term stochastic makes a difference in countries like Sweden or France. This is not possible using the same software that I’m using. But it is easy to estimate with other software. I suspect that the drift term will have almost no stochastic variation, because the time series already look very linear.
3. Figure out how to get more meaningful estimates of n_t for the United States and other countries where estimation of the latent k_t doesn’t produce much of interest. One way to do this might be to change the assumptions about the statistical distribution of n_t . For example, it could have heavier tails (which is one way to make it more sporadic). Doing this requires other software, but I think I use a Bayesian approach for estimating all of the models using a consistent method.

4. It might be interesting to include more countries and show how shocks are correlated. (In Asia, we can add Taiwan, Korea and HK). In Europe there are many more countries. And we can add NZ and Australia.
5. It might be interesting to do a kind of “coherent” forecast in which correlation of the $n(t)$ by country are taken into account. Perhaps there is also correlation of the $\epsilon(t)$?
6. It would be interesting to see how much difference it makes to forecast with this latent model rather than the observed k_t . We could take a country like France or Sweden where there appears to be a difference and show how the forecast would have evolved over time, particularly for launch years which show a large gap between the observed and the latent $k(t)$ values.